


```

# Test different values of k
k_values = range(1, 20)
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train, y_train, cv=5)
    print(f'k={k}, Mean Accuracy: {np.mean(scores)}')

def lazy_classification(X_train, y_train, X_test, y_test, k_value):
    # Train k-NN with the selected k
    knn = KNeighborsClassifier(n_neighbors=k_value)
    knn.fit(X_train, y_train)

    # Predictions on the test set
    y_pred = knn.predict(X_test)

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    print("Confusion Matrix:")
    print(cm)

    # Classification Report
    cr = classification_report(y_test, y_pred)
    print("\nClassification Report:")
    print(cr)

    # Evaluate the model on the test set
    accuracy = knn.score(X_test, y_test)
    print(f'Test Set Accuracy: {accuracy}')

# Example usage
tune_knn_parameters(X, y)

```

```

k=1, Mean Accuracy: 0.8998101265822784
k=2, Mean Accuracy: 0.8972784810126584
k=3, Mean Accuracy: 0.9148417721518987
k=4, Mean Accuracy: 0.9123734177215189
k=5, Mean Accuracy: 0.919873417721519
k=6, Mean Accuracy: 0.9223101265822784
k=7, Mean Accuracy: 0.9248417721518987
k=8, Mean Accuracy: 0.9298101265822785
k=9, Mean Accuracy: 0.9323417721518987
k=10, Mean Accuracy: 0.9298417721518988
k=11, Mean Accuracy: 0.9298417721518988
k=12, Mean Accuracy: 0.9298417721518988
k=13, Mean Accuracy: 0.9298417721518988
k=14, Mean Accuracy: 0.9223101265822784
k=15, Mean Accuracy: 0.9273101265822785
k=16, Mean Accuracy: 0.9248101265822786
k=17, Mean Accuracy: 0.9273101265822785
k=18, Mean Accuracy: 0.9248101265822786
k=19, Mean Accuracy: 0.9198101265822786

```

In [76]: `lazy_classification(X_train, y_train, X_test, y_test, 18)`

Confusion Matrix:

```
[[58  2]
 [ 8 32]]
```

Classification Report:

	precision	recall	f1-score	support
B	0.88	0.97	0.92	60
M	0.94	0.80	0.86	40
accuracy			0.90	100
macro avg	0.91	0.88	0.89	100
weighted avg	0.90	0.90	0.90	100

Test Set Accuracy: 0.9

In [77]: `import numpy as np`

```

def plot_learning_curve(estimator, title, X, y, cv=5, n_jobs=-1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Parameters:
    - estimator: The entity employed to fit the data.
    - title: The chart's title.
    - X: Training vector, where n_samples represents the number of samples, and n_features signifies the number of
    - y: Target corresponding to X for classification or regression.
    - cv: Cross-validation generator or an iterable, default=None.
    - n_jobs: The number of parallel jobs to execute, default=None.
    - train_sizes: Relative or absolute quantities of training examples utilized for generating the learning curve.
    """

```

Returns:

- plt: Matplotlib plot object.

```
"""
plt.figure()
plt.title(title)
plt.xlabel("Training examples")
plt.ylabel("Score")

train_sizes, train_scores, test_scores, fit_times, _ = \
    learning_curve(estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes, return_times=True)

train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

plt.grid()

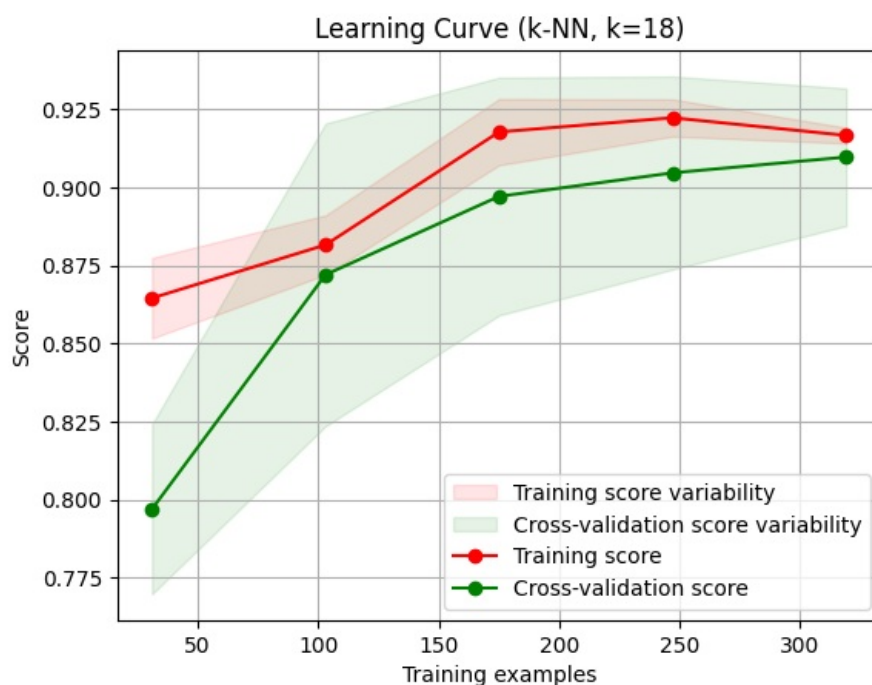
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r", label="Training score variability")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1,
                 color="g", label="Cross-validation score variability")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

# Assuming X_train, y_train, X_test, y_test are your training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Choose a k value for k-NN
k_value = 18
knn = KNeighborsClassifier(n_neighbors=k_value)

# Plot learning curve
plot_learning_curve(knn, f"Learning Curve (k-NN, k={k_value})", X_train, y_train, cv=5, n_jobs=-1)
plt.show()
```



```
In [78]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Assuming X, y are your dataset
# Replace this part with your actual data loading and preprocessing
# For example:
# df_heart_disease = pd.read_csv("your_dataset.csv")
# X = df_heart_disease.drop('target', axis=1)
# y = df_heart_disease['target']
```

```

# X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Function for tuning k-NN parameters
def tune_knn_parameters(X, y, k_value):
    # Split the data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    # Train k-NN with the selected k
    knn = KNeighborsClassifier(n_neighbors=k_value)
    knn.fit(X_train, y_train)

    # Convert 'B' to 0 and 'M' to 1 in y_test
    y_test_binary = y_test.map({'B': 0, 'M': 1})


    # Plot ROC curve
    plot_roc_curve(knn, X_test, y_test_binary, f"ROC Curve (k-NN, k={k_value})")



# Function to plot ROC curve
def plot_roc_curve(model, X_test, y_test, title):
    y_prob = model.predict_proba(X_test)[:, 1]
    fpr, tpr, thresholds = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)



    plt.figure()
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = {:.2f})'.format(roc_auc))
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(title)
    plt.legend(loc="lower right")
    plt.show()

# Assuming k_values is the list of k values you want to try
# For example, k_values = [1, 3, 5, 7, 9]
for k_value in [1, 3, 5, 7, 9]:
    tune_knn_parameters(X, y, k_value)

```







```
In [79]: # Using Decision Tree Classifier from sklearn.tree import DecisionTreeClassifier

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=40)

# Decision Tree Classifier
# Tuning the max depth parameter
best_depth = None
best_accuracy_tree = 0

# Add more values for tuning max depth
for depth_value in [None, 5, 10, 15]:
    # Train Decision Tree Classifier with the current max depth
    dtc = DecisionTreeClassifier(max_depth=depth_value, random_state=42)
    scores = cross_val_score(dtc, X_train, y_train, cv=5, scoring='accuracy')

    # Compute mean accuracy
```



```

mean_accuracy_tree = scores.mean()

print(f"Max Depth={depth_value}, Mean Accuracy: {mean_accuracy_tree}")

# Update best parameters if needed
if mean_accuracy_tree > best_accuracy_tree:
    best_accuracy_tree = mean_accuracy_tree
    best_depth = depth_value

# Choose the optimal max depth based on the tuning results
optimal_depth = best_depth

# Perform classification with the optimal max depth
dtc = DecisionTreeClassifier(max_depth=optimal_depth, random_state=42)
dtc.fit(X_train, y_train)

# Predictions on the test set
y_pred_tree = dtc.predict(X_test)

# Decision Tree Classifier - Confusion Matrix
cm_tree = confusion_matrix(y_test, y_pred_tree)
print("\nDecision Tree Classifier - Confusion Matrix:")
print(cm_tree)

# Decision Tree Classifier - Classification Report
cr_tree = classification_report(y_test, y_pred_tree)
print("\nDecision Tree Classifier - Classification Report:")
print(cr_tree)

# Decision Tree Classifier - Evaluate the model on the test set
accuracy_tree = dtc.score(X_test, y_test)
print(f'Decision Tree Classifier - Test Set Accuracy: {accuracy_tree}')

```

Max Depth=None, Mean Accuracy: 0.934873417721519
Max Depth=5, Mean Accuracy: 0.9298417721518988
Max Depth=10, Mean Accuracy: 0.934873417721519
Max Depth=15, Mean Accuracy: 0.934873417721519

Decision Tree Classifier - Confusion Matrix:
[[57 3]
[1 39]]

Decision Tree Classifier - Classification Report:

	precision	recall	f1-score	support
B	0.98	0.95	0.97	60
M	0.93	0.97	0.95	40
accuracy			0.96	100
macro avg	0.96	0.96	0.96	100
weighted avg	0.96	0.96	0.96	100

Decision Tree Classifier - Test Set Accuracy: 0.96

```


In [80]: # Assuming `model` is your trained model and `X`, `y` are your features and labels
train_sizes, train_scores, test_scores = learning_curve(dtc, X, y, cv=5)

# Calculate mean and standard deviation across cross-validation folds
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

# Plot learning curve
plt.figure(figsize=(10, 6))
plt.fill_between(train_sizes, train_scores_mean - train_scores_std, train_scores_mean + train_scores_std, alpha=0.1)
plt.fill_between(train_sizes, test_scores_mean - test_scores_std, test_scores_mean + test_scores_std, alpha=0.1)
plt.plot(train_sizes, train_scores_mean, 'o-', color="r", label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g", label="Cross-validation score")

plt.title("Learning Curve")
plt.xlabel("Training Examples")
plt.ylabel("Score")
plt.legend(loc="best")
plt.show()

```



Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js